Perception of Virtual Characters

Eduard Zell
Trinity College Dublin

Katja Zibrek
Trinity College Dublin

Rachel McDonnell
Trinity College Dublin

ABSTRACT

This course will introduce students, researchers and digital artists to the recent results in perceptual research on virtual characters. It covers both how technical and artistic aspects that constitute the appearance of a virtual character influence human perception. We will report results of studies that addressed the influence of low-level cues like facial proportions, shading or level of detail and higher-level cues such as behavior or artistic stylization. We will place emphasis on aspects that are encountered during character development, animation, and achieving consistency between the visuals and storytelling. The insights that we present in this course will serve as an additional toolset to anticipate the effect of certain design decisions and to create more convincing characters, especially in the case where budgets or time are limited.

KEYWORDS

perception, virtual character

ACM Reference Format:
Eduard Zell, Katja Zibrek, and Rachel McDonnell. 2019. Perception of Virtual Characters. In *Proceedings of ACM Conference (SIGGRAPH Courses'19).* ACM, New York, NY, USA, 17 pages. https://doi.org/10.1145/nnnnnnn.nnnnnnn

1 COURSE DESCRIPTION

Virtual humans are finding a growing number of applications, such as in social media apps, Spaces by Facebook, Bitmoji and Genies, as well as computer games and human-computer interfaces. Their use today has also extended from the typical on-screen display applications to immersive and collaborative environments (VR/AR/MR). At the same time, we are also witnessing significant improvements in real-time performance and increased visual fidelity of characters. The question of how these developments will be received from the user's point of view, or which aspects of virtual characters influence the user more, has therefore never been so important. This course will provide an overview of existing perceptual studies related to the topic of virtual characters.

To make the course easier to follow, we start with a brief overview of human perception and how perceptual studies are conducted in terms of methods and experiment design. With knowledge of the methods, we continue with artistic and technical aspects which influence the design of character appearance (lighting and shading, facial feature placement, stylization, etc.). Important questions on character design will be addressed such as - if I want my character to be highly appealing, should I render with realistic or stylized shading? What facial features make my character appear more trustworthy? Do dark shadows enhance the emotion my character is portraying?

Finally, we dive deeper into the movement of the characters, exploring which information is present in the motion cues and how motion can, in combination with character appearance, guide our perception and even be a foundation of biased perception (stereotypes). Some examples of questions that we will address are - if I want my character to appear extroverted, what movement or appearance is needed to achieve this? Can character appearance influence my moral decisions in a video game? Does my behaviour towards a virtual human change depending on my empathy level towards them?

The course provides the overview of the relevant studies in a way that makes it easy to identify answers to practical questions in production and character development. At the same time, we avoid giving definite answers to questions of character design and encourage further investigation by listing questions left unanswered to allow for critical evaluation of the presented research.

Finally, participating in a perceptual experiment is a multi-modal experience, which cannot be reproduced only by descriptive reports of the experiment design. For this reason, we will select a few representative experiments and run a highly compact version of them during the course for illustration purposes. The stimuli will be shown on the projector wall and the participants will be able to rate the stimuli within a small time-frame using their smartphones and QR codes. Experiments will primarily be selected to introduce a new topic. We are fully aware that the obtained results are not representative by any means, but we believe that such live surveys will improve the understanding of the study design, increase engagement of participants, and be a welcoming break during a 90 minute talk.

Related Tutorials and Courses. Courses of the last 10–15 years at SIGGRAPH, SIGGRAPH Asia or Eurographics covered topics such as experiment design [Cunningham and Wallraven 2013], visual perception of simple 3D shapes [Fleming and Singh 2009], as well as perception in graphics with applications to display technologies and virtual environments [Glencross et al. 2006a; Mania and Reinhard 2008; Thalmann et al. 2007]. Other courses covered a mixture of low-level stimuli perception and application in graphics, where character perception was partly addressed as well [McNamara et al. 2011; O'Sullivan et al. 2004]. Finally, there are courses that focused on perception of specific aspects of virtual characters; these include: (i) the expressiveness of body motion [Hertzmann et al. 2009], (ii) crowds [Badler et al. 2014; Donikian et al. 2009; Huerre et al. 2010; Thalmann et al. 2007] (iii) a multidisciplinary study of emotions covering aspects of philosophy, psychology and physiology [Geslin 2012] and (iv) the creation of believable characters for dialogues [Jung et al. 2011].

SIGGRAPH Courses'19, July 2019, Los Angeles, USA

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM…$15.00
https://doi.org/10.1145/nnnnnnn.nnnnnnn

Our course is the first to cover perception of virtual humans in a single resource, and addresses much more recent work than previous courses. We feel that it will be accessible for non-experts and a starting point for further investigation on related topics.

Audience. This course is suited for students, who want to get an overview of recent developments of perceptual research on virtual characters and identify open topics. Furthermore, this course is particularly designed for researchers and artists who work on virtual characters but are less familiar with the perceptual research.

Prerequisites. Fundamentals about creating and animating virtual characters and knowledge about design and analysis of perception experiments is beneficial, but not required.

Difficulty. Beginner to Intermediate

Website. www.eduardzell.com/VirtualCharacters.html

2 SCHEDULE

I **Visual Perception Basics** (20 minutes)
- (a) Perception
- (b) Experiment Design & Statistics
- (c) Stimuli Creation

II **Character Appearance** (35 minutes)
- (a) Character Stylization
- (b) Character Realism
- (c) Facial Proportions
- (d) Level of Detail
- (e) Skin Appearance
- (f) Lighting and Shading
- (g) Visual Attention of Facial and Body Parts

III **Character Motion and Behaviour** (35 minutes)
- (a) Emotion
- (b) Gender
- (c) Gender and Emotion Bias
- (d) Personality

3 LECTURER BIOGRAPHIES

Eduard Zell

Eduard Zell is a postdoctoral Research Fellow at Trinity College Dublin. In 2018, he received his PhD from Bielefeld University, Germany on the topic of creation, animation and perception of virtual faces. His work was published at SIGGRAPH and Scientific Reports and for his thesis, he received the best thesis of the faculty award as well as the Eurographics PhD Award. Prior to his PhD, he completed a highly practical degree in Computer Animation and Visual Effects (M.Sc.) at Bournemouth University, UK.

Katja Zibrek

Katja Zibrek is a postdoctoral Research Fellow at Trinity College Dublin (TCD). She holds a diploma in Psychology (University of Ljubljana, Slovenia) and a PhD in Computer Science (TCD). She has conducted research in the area of perception in graphics, particularly in investigation of gender, emotion an personality perception of virtual characters. She has published at SIGGRAPH, ACM Transactions on Applied Perception and IEEE Transactions in Visualisation and Computer Graphics.

Rachel McDonnell

Rachel McDonnell is an Assistant Professor at Trinity College Dublin. Her research interests include perception, animation, and virtual humans. She has been a member of many IPCs, including the Eurographics and SIGGRAPH papers committee and has published over 50 papers on topics in perception, facial animation, and virtual humans. She has served as both program and conference chair of the Symposium on Applied Perception, and is on the editorial board for the associated journal - Transactions on Applied Perception.

4 VISUAL PERCEPTION BASICS

Perception is an important part of graphics, and the knowledge of how the human visual system interprets visual stimuli was the foundation of many techniques used in graphics. For example, the knowledge that people can perceive a sequence of images, depicting a moving object, when presented in fast succession as fluid motion, was a foundation for animation. Another example is that people see color due to only three types of light sensitive receptors or "cones" on the retina (red, blue and green). This was the basis of the RGB color system, where any color could be expressed as an integer on the spectrum of red, green and blue. There are many more examples of how graphics exploits these basic traits of the human visual system, however, this course will only make a brief overview of them. The reader is encouraged to refer to the previous courses [Fleming and Singh 2009; Glencross et al. 2006b] which cover basic visual perception of 3D environments in more detail. Here, we present a broader definition of perception, which is needed to understand some concepts of virtual character perception.

4.1 Perception

Perception comes from the Latin word perception which literally means "to seize" or "to understand". In order to perceive the world, stimuli need to access the organism through a system, developed to turn information into the activity in the nervous system – a process called "sensation" (visual, tactile, audio, etc.). However, these sensations need to be organized in a meaningful experience. This part is mainly performed by the brain, which processes information from the senses and interprets its relevance to the organism. The first process of transforming sensory–driven information, is also called the **bottom–up** process. The second process is based on acquired information about the world through learning and provides a context for the information from the senses to be interpreted, also known as the **top–down** process. A simple model of the interplay of the processes is shown in Figure 1.

Bottom–up process transforms low–level sensory information into high–level information. This is needed because visual stimuli are very complex. A good example of this process is depth perception – the brain needs to see depth from images, which are displayed in an eye's retina as two dimensional images. So how does the brain do it? Well first, it uses information from two eyes, which deliver a slightly different angle of the image (retinal disparity). Then, it joins the images in the visual cortex to assess the depth information. To do so accurately and fast, it uses a set of learned or predefined rules of organization.

These are usually known as principles of visual organization (also Gestalt principles) which are part of the top-down process. A common principle is "figure and ground" (Rubin [1915], described in [Beardslee and Wertheimer 1958]) where the figure is seen as a meaningful object in the field of view (typically presented as a smaller, connected image), and the ground is the less relevant background. "Grouping" is another form of organization, where our perceptual system joins separate objects together to create a whole by their visual proximity, similarity, continuity, closure, symmetry and common fate [Wertheimer 1923]. Very important principles regard depth perception. These principles are especially important in art and graphics. For example, to simulate depth in 2D images, (Figure 2a) we can use two lines, converging into a point, which create the illusion that they are actually parallel and continue into a distance. The two yellow lines in the picture have the same length but because they are put on the "path" of the two converging lines they appear to be at different distances, therefore creating the illusion that one is longer than the other. A lot of principles of organization have been observed with these types of visual illusions. The illusions signify that there are competing processes happening in visual perception – for example, in the figure-ground principle, if the visual process cannot determine what is figure and what ground, it will switch between them, depending on the person's attention focus (Figure 2b).

While it is not precisely known which of these principles of visual organization are learned through experience with the world or which are biologically inherited, many still belong to a category of bottom–up perception and are processed in the primary visual cortex, the brain region which processes the most fundamental visual stimuli, such as orientation and color [Schwarzkopf et al. 2011]. However, it is known that our top–down processes, such as attention, expectations, motivations, etc., influence our perception as well. These mental representations or "schemas" include everything we learned about the world and provide a fast assessment of the meaning of the stimuli, especially when sensory information is vague or ambiguous. For example, top–down perception is the reason why we can perceive a human form from a simple point-light display [Johansson 1973], due to our previous experience with observing people. However, this ability is also the reason for some erroneous judgments. A practical example is the failure to notice spelling mistakes in the text since we can derive meaning from words even when we put attention on a few letters of those words. On the other hand, these failures can be used to optimize graphics content without introducing a perceptible change (see for example mesh simplifications based on a perceptually driven technique [Gibson and Hubbold 1997]).

When perceiving virtual characters, the understanding of both bottom-up and top-down processes are important. Knowing how the 3D shape and depth can be perceived from a 2D representation is just as crucial for character design as understanding the importance of pre-existing schemas when observing virtual humans. These schemas can be very broad and, as we will discuss in the last section of this course, include social perception. In social situations, people frequently make judgments of other people based on very little information. Visual appearance can result in attributing particular personality traits even to strangers, which results in perceiving any behavior from that person in the context of the first judgment. Stereotypes are such an example, where other peoples' behaviour is analyzed according to a group or category they belong to (e.g., race, gender, age) and not unique constellations of their personal attributes. In graphics, we can also use bias as a measure of authenticity of a designed virtual character - a virtual human who is perceived to be realistic in behaviour and appearance, could also induce a biased response [Zibrek et al. 2018].

The recognition of these top–down perceptual effects is also important when designing a perceptual experiment. Participants themselves come from various backgrounds and are in different mental states at the time of the experiment, which might interfere

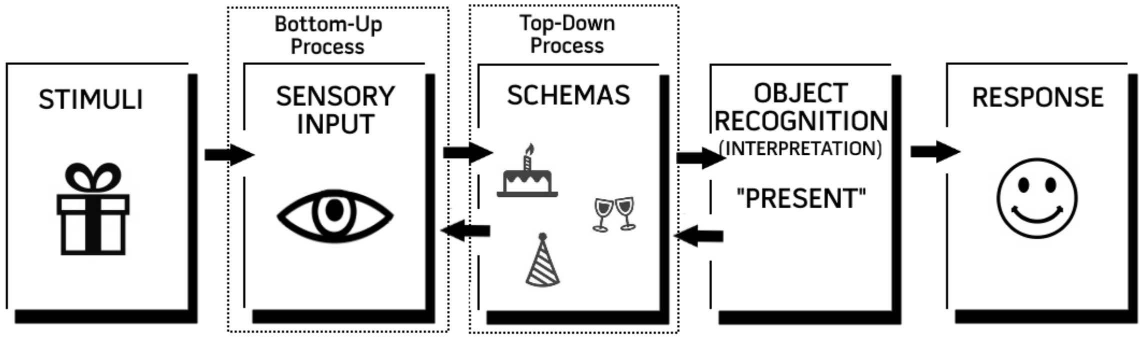

Figure 1: Simple model of perception. The object is first seen by the sensory system (visual input) and then interpreted in context (being at a birthday party). The object is recognized (present) and ends in the response of happiness. The arrows represent the direction of effect, where there is a feedback loop between sensory information (bottom–up) and already acquired representations of the world (top–down). Top–down processes also include cognitive states of the person (his attention, motivation, personality) which can affect the sensory organ to perceive selectively. For example, if the person does not trust his friends that gave him a present, he might react concerned or even frightened at the sight of the box. There is also an example of perceptual organization (grouping) in this image: the boxes that describe separate steps in the model are registered by the sensory system as broken lines yet we perceive them a rectangular objects, seemingly casting a shadow.

Figure 2: Examples of visual illusions: a) Ponzo illusion, where the yellow lines appear to be different even though they are of the same length; b) Rubin vase, where the object in the image can be seen as a vase or two faces.

with the variables we are measuring. The following section provides a short overview of the basics of experiment design for the investigation of the perception of virtual characters.

4.2 Experiment Design & Statistics

There are some specifics of the experimental design when investigating the perception of virtual characters. In this section, we are going to cover the basics by introducing the following topics:

- The influence of participants (sampling)
- The differences in measures: direct and indirect
- An overview of the most important concepts, when analyzing the data (e.g., mean, standard deviation, trust interval and the difference between significant and intermediate results).
- The importance of controlled stimuli

4.2.1 Participants. As mentioned in the previous section on perception basics, participants come from a variety of backgrounds and this variation can introduce unwanted effects on our collected data. Increasing the number of participants is usually the best way to avoid any effect of individual variation, as this gets dispersed throughout a large sample while only the systematic effects should remain. Another important key is also in the sampling approach - in order to generalize results from the sample of the population, which is the basic premise of inferential statistics used in perceptual studies, the sample should represent the population well. If, for example, we got a result that the recognition of emotion is higher for stylized than realistic characters, and our sample is 30 male and 5 female participants, the conclusion cannot be generalized to all people, only males. It is similar with sampling only from the university campuses, or collecting data from only one cultural group.

In reality, getting a perfect sample for our experiment is usually extremely challenging. However, some improvements can be made simply by insuring that the selection of participants is as randomized as possible. Another way is to include previous knowledge on the perception of people and run pilots or pre-tests to test the experiment design. For example, when investigating recognition of emotion of virtual characters, research on the psychology of emotions shows that the perceivers' own emotions influence emotion perception of others [Niedenthal et al. 2000]. If we are interested to control for this effect, we could measure the emotional state of the participant prior to the experiment and include it in the statistical analysis.

4.2.2 Measures. Peoples' attitudes towards virtual characters can be measured in various different ways. The most commonly used measures are subjective responses, where people are asked to give answers to a questionnaire, such as rate their experience or make a decision about what they had witnessed. Subjective responses are usually obtained by questionnaires, where Likert scales and semantic differential scales are used in the attempt to quantify data. Likert scale prompts the person to give a rating of an agreement with a particular statement (e.g., "On a scale from 1 to 7, how eerie is the character?"), while the semantic differential scale has two different descriptors on each end of the scale, for which it was previously established that they belong to the same dimension (e.g., an emotional response scale can range from happy to sad). These scales have certain disadvantages. People's intention, mood, personality type and other unrelated factors can influence the way they give answers - some people avoid giving extreme ratings, develop ideas and strategies on how to assess stimuli, give intentionally misleading answers, etc. [Bartz et al. 2008]. In order to control for this, repetitions of the same question can give a more reliable result, and there are even tests which measure people's willingness to give socially desirable answers (Social Desirability Scale, Crowne and Marlowe [1960]). The most common approach, however, is to use standardized tests, which are created from a set of scales, measuring a specific construct, and have been tested on a large sample and controlled for validity (that the test is measuring the intended construct) and reliability (the test measures the construct consistently across time, individuals and situations). An example of a standardized measure which measures attitudes towards artificial humans is the Godspeed Questionnaire, introduced by Bartneck et al. [2009] and revised by Ho et al. [2010]. This instrument uses 4 indices with high internal reliability - warmth, humanness, eeriness and attractiveness. A lot of research studies, however, do not use the same terminology and the lack of universality is a known issue in the field of character perception [Kätsyri et al. 2015].

Another way to avoid subjective mapping of answers, a forced choice task can be used, where a limited range of options is given, and the participant must choose the one which is the closest to his answer. In the Two-Alternative Forced Choice (2AFC) experiment design, speed and accuracy of choices between two alternatives given a timed interval are tested [Blackwell 1952]. Most of the low level perceptual experiments use some version of this task, and the goal is to retrieve the thresholds of stimuli detection or the levels of when the stimuli changes the perceived intensity. An example of how this test could be used for evaluating the perception of virtual characters, is in virtual crowds or the so called detection of 'imposters' [Hamill et al. 2005], where simplified versions of characters are introduced to increase the rendering speed of a large crowd without being noticed by the viewer. An extended version of the 2AFC measure is the multi-dimensional scaling method (MDS), where viewers do not only report the detection of change but also the degree to which the stimuli changes (see for example Logvinenko and Maloney [2006]).

Brain studies using fMRI and EEG can be used as well. Experimenters can monitor participant's heart rate, respiratory rate and skin conductance to track changes in anxiety levels of people who are observing the character. Peoples' eye gaze can give a lot of information on their attention to particular areas of the character, indicating areas of interest or disturbance. These objective measures are referred to as physiological measures and have many advantages: they are quantifiable and do not require participants' conscious evaluation. Reasons against using these measures could be poor accessibility and cost of the machines, additional expertise for analyzing the results and non-direct association between physiological and mental responses.

Indirect measures are therefore the ones where the participant is not aware of the purpose of the testing and cannot affect the outcome by conscious processing. For example, rather than asking the participant how threatening the character appears to him or her, we can measure participant's increase in heart rate. Other indirect measures are based on semantic priming (e.g., Stroop test [Jensen 1965], Implicit Association test [Greenwald et al. 1998]) and are also used to study the perception of virtual characters [Banakou et al. 2013]). The later study also utilizes virtual reality as a tool, where ecological validity of behaviour is possible (e.g., the measure of proximity to virtual humans in the studies of Bailenson et al. [2005]). Indirect measures are extremely valuable as they bypass any conscious interpretation from the person which could affect the measured data. However, indirect measures may pose a question to validity - do these measures really reflect the nature of the studied construct? To increase validity, a combination of direct and indirect measures is usually the best choice for a rigorous perceptual study.

Peoples' responses can also be collected through observation and qualitative measures. These methods are helpful when we do not have much knowledge about a particular problem we are investigating and do not know how to approach it. A Q qualitative approach will provide a wide range of data but it will be difficult to analyze in a concise way and subject to noise in the data.

4.3 Stimuli Creation

Obtaining images or videos of adequate quality and that fit the purpose of a perceptual experiment may be a difficult task. In general one is either interested in stimuli that change concisely one single aspect or in a large collection of stimuli, such that inconsistencies will vanish later as variance within the statistical analysis. In the following we list several methods to obtain stimuli of virtual characters together with the advantages and pitfalls of each stimuli creation method.

- Collecting images from the internet may be tempting and is certainly the easiest way to obtain big databases. The downside is that many aspects (backgrounds, light, dresses, image resolution and aspect ratio etc.) cannot be controlled and stimuli taken from blockbuster productions cannot be added to the submission without separate copyright agreements with the publishers, which are difficult to obtain. The more specific the requirements for the stimuli, the more difficult it becomes to find the right images. In some cases consistency can be improved by post-processing the selected images (e.g., replacing the background or color correction).
- Morphing between images or 3d models has the advantage that it can be easily accomplished with specialized software. It is also the fastest method to achieve very fine-grained sampling between a photograph and a virtual character. The downside of this approach is that the interpolation is defined by technical and not artistic terms. Visual artifacts may arise if the images are highly different.
- Creating stimuli manually by a visual artist offers a high level of control on the final result at the cost of being very time-consuming. Caution must be paid to subjective traits that should be ideally cross-evaluated to prove that the artist achieved the intended goal.
- Many properties can be controlled in computer graphics by modifying a few parameters. Such parametrizations greatly

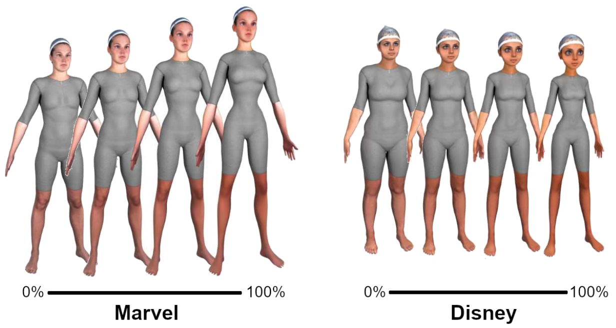

Figure 3: Transitions between virtual characters from realistic to specific stylizations, sampled by 0%, 33%, 66% and 100% [Fleming et al. 2016].

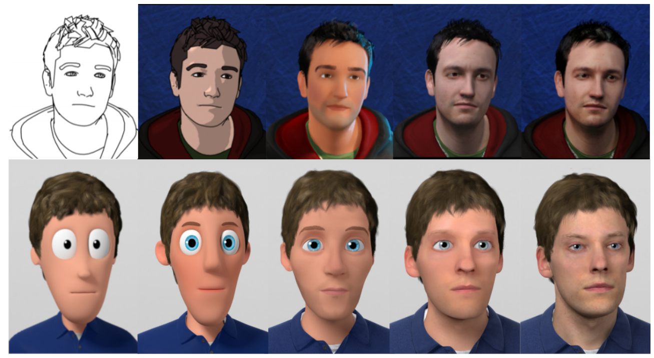

Figure 4: Examples of consistent stylization levels, starting from most abstract to most realistic. *Top:* stylization of rendering styles [McDonnell et al. 2012]. *Bottom:* stylization of shape and material [Zell et al. 2015].

simplifies creation of highly consistent stimuli (e.g., controlling the light intensity). However, it is unlikely that an equal sampling of the parameter space will also create an equal sampling of the final appearance. The relationship of parameters and the final appearance is in general non-linear.

- For transferring one property (e.g., texture, motion) from one character to another, a mapping or cross-parametrization algorithm will be required. Such algorithms are designed with certain implicit assumptions that may or may not apply to the specific case. For example, algorithms that align 3d meshes try to keep area distortions small. While obvious limitations of the algorithm can be fixed manually, there is also a chance that the algorithm itself introduced unintended effects.

5 CHARACTER APPEARANCE

In this section, we will focus on visual characteristics of virtual characters that are largely not affected by temporal changes. Some topics like Level of Detail (Section 5.4) or Lighting and Shading (Section 5.6) have a rich history of general perceptual studies that have only a partial overlap with character related topics. In such cases only a selection of relevant studies will be discussed.

5.1 Character Stylization

Virtual characters are highly diverse, making comparisons between each other a difficult task. Even in terms of categorization, no unique definition exists. Based on the observation of different levels of abstraction in comics, McCloud [1993] classifies stylization along the iconic and non-iconic scales. A realistic face becomes a smiley under iconic stylization, or a cubist portrait under non-iconic abstraction. Ritchie et al. [2005] extends this concepts by introducing hyper-realistic characters, a category for characters like The Hulk and Golumn, who look highly realistic but do not exist in real life.

Similar concepts exist for perceptual studies, however a stronger focus is put on stimuli consistency. The two most common scales are photo-realistic vs. stylized/iconic and photo-realistic vs. anthropomorphic or hyper-realistic. It is also common to subdivide the realistic vs. abstract scale further by subdividing the abstraction level into technical components like the general form (shape), resolution of a mesh (tessellation), surface properties (shading, albedo texture etc.) and finally motion. Such subdivision reflects both the different styles in artwork and the technical limitations. Furthermore, the fine grained subdivision allows to track back the contribution and importance of the different ingredients. The downside of this approach is that testing all these parameters becomes difficult, especially because the combinations of parameters grow exponentially. For example testing five characters with four levels of stylization and two different shaders, requires 40 different combinations. If we just increase each scale by one additional sample, the number of stimuli increases to 90. From a practical point of view, creating that many combinations might not even make sense as some combinations would never be used in practice. However, in terms of data analysis, equal sampling is preferred.

Several studies investigated how virtual characters are perceived in terms of realism or appeal across different stylization levels. However, stimuli in early work were mainly based on pictures taken from commercial productions [Dill et al. 2012; Hanson 2005; MacDorman 2006; Schneider et al. 2007]. The lack of consistency (different characters, lighting, backgrounds etc.) of such stimuli, do not always guarantee that the origin of the measured effect is only caused by the different stylizations. For this reason, we focus in the following on studies with more controlled stimuli sets.

Body Perception. Fleming et al. [2016] evaluated the appeal and realism of female body shapes, which were created as morphs between a realistic character and stylized versions following design principles of major computer animation studios (Figure 3). Surprisingly, the most appealing characters were in-between morphs, where 33% morphs had the highest scores for realism and appeal and 66% morphs were rated as equally appealing, but less realistic.

Faces. Wallraven et al. [2007] studied the perceived realism, recognition, sincerity, and aesthetics of real and computer-generated facial expressions using 2D filters to provide brush, cartoon, and illustration styles. They concluded that realistic depictions improve subjective certainty about the conveyed expression. Later, they evaluated the perceptual realism of computer-generated faces under progressively blurred normal vectors and textures, finding no effect with their setup [Wallraven et al. 2008]. In the study by McDonnell et al. [2012] the authors investigated the impact of different rendering styles on the appeal and trustworthiness of the characters (Figure 4, top). In contrast to most studies, this was done for static renderings and short animations. Rendering styles that were close to the most basic shading model in computer graphics were rated as less appealing and trustworthy. Motion amplified this effect. By separating the stylization across shape and material independently, Zell et al. [2015] (Figure 4, bottom) identified that: (i) Shape is the main descriptor for realism, and material increases realism only in case of realistic shapes. (ii) Strong mismatches in stylization between material and shape negatively affect the appeal and attractiveness of the characters and make them eerier. (iii) The albedo texture modifies primarily the perceived changes of the material and blurring a realistic textures a make-up effect can be achieved - the character appeal and attractiveness increases, without reducing realism. (iv) Finally, abstract characters with realistic materials were perceived as highly eerie, validating the design choices of some horror movies with living puppets.

Instructors. One rather frequently encountered use-case of virtual characters are instructors or experts. Despite strong progress over the last years, virtual characters can still be reliably detected in most cases [Mader et al. 2017]. This raises the question whether they are perceived as competent as real humans. Confronted with an ethical dilemma decision within a medical context, participants had to make a choice before and after the advise of a doctor (expert). The doctor in the video sequence was either a virtual avatar or a human. The recommendation of the doctor had a significant influence on the decision of the participants independently of his appearance. This was even the case when the motion was modified to be "jerky". Similar results were obtained later in a follow-up study [Dai and MacDorman 2018]. However, not all studies come to the same conclusions. Testing the learning outcome of a recorded lecture with slides enriched with a small video of the real instructor, a virtual avatar or a robot, the learning outcome varied. It was smallest for the virtual avatar. Interestingly, students liked the virtual avatar as much as the real person and disliked the robot [Li et al. 2016].

Figure 5: A real person together with a digital double in a virtual environment (CAVE) [Waltemate et al. 2018].

Another study tested whether stylization level had an influence on expert identification and in consequence whether trust is influenced by stylization [Pan and Steed 2016]. Within the study, participants had to answer several difficult question and were assisted by two personalities, where only one was an expert. The personalities were either digital avatars, humans or a humanoid robot. While the robot was placed in front of the wall, the two other were projected on a wall. If the digital avatar was the expert, participants struggled to identify him. In contrast, experts represented by the robot or by a real person were identified reliably. Finally, a meta-analysis comparing the subjectively and objectively measured benefit of adding human-looking virtual avatars as an interface comes to the conclusion that adding an avatar is beneficial, but the effect size is small [Yee et al. 2007].

5.2 Character Realism

With computer graphics reaching closer and closer an indistinguishable level of photo-realism, questions remain on how close we are to this goal and how do we react towards virtual doubles of ourselves.

Becoming Real. Earlier studies focused explicitly on the identification of the boundary when characters are perceived as real, by morphing between photographs and puppet faces [Looser and Wheatley 2010] or between photographs and virtual faces [Cheetham et al. 2011]. The results of these works indicate that this question is indeed a categorical decision and that characters must match a high level of realism until they are perceived as real. Furthermore, it seems that eyes and mouth contain the most relevant information followed by the nose, while skin is less relevant (see Section 5.7). Interestingly, due to the visual quality of virtual doubles and output from machine learning algorithms, the identification of computer generated characters is gaining interest within forensic research. Recent studies show that participants exposed to training, feedback within the trial, and incentives were able to classify up to $85-90\%$ of images correctly as real or computer generated compared to a performance of $\approx 50\%$ in mechanical Turk experiments without additional incentives [Mader et al. 2017]. In terms of error detection, shading is more important than color and compared to other facial areas judging by the eyes alone reveals the highest accuracy [Fan et al. 2014].

Doppelgänger. Recent developments in 3d scanning facilitate new types of experiments where a realistic virtual double of participants is created during the experiment (Figure 5). Such virtual doubles or doppelgängers allow testing the importance of having a virtual character that are either a look-alike or non-similar to users. Within an interactive application [Fox and Bailenson 2009], participants could control the weight of a virtual character by doing physical exercise. If the virtual character was a doppelgänger, participants exercised significantly more. Furthermore, advertisement with doppelgängers tends to be more effective [Ahn and Bailenson 2011] and participants react less aggressive to doppelgängers of others in games [Segovia and Bailenson 2012]. However, watching at a virtual double is not always of benefit. In preparation of public speaking, participants who watched non-similar characters reduced anxiety compared to watching virtual doubles giving the

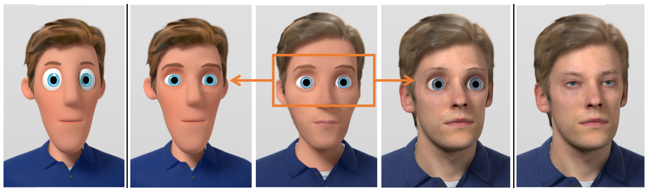

Figure 6: Illustrative example for changing facial proportions across different stylization levels. Eyes of the character in the middle are transferred to abstract (*left*) and realistic (*right*) characters. While the abstract character still look reasonable, this is not the case for the realistic character.

talk [Aymerich-Franch et al. 2014]. Virtual doubles, in combination with motion capturing increase body ownership and presence within the virtual environment and facilitate the acceptance of the virtual body as their own [Gorisse et al. 2019; Waltemate et al. 2018]. Finally, accuracy of body weight estimation is independent of the participant's gender [Thaler et al. 2018]. However, males accepted a larger weight range as their own. Females but not males considered a thinner body as ideal.

5.3 Facial Proportions

In Seyama and Nagayama [2007] stimuli were created by morphing between photographs of real people and dolls. By controlling individually the morphing speed of facial parts as well as by scaling facial parts individually, it was found that realistic characters were perceived as less appealing if facial parts had strong deviations in terms of size (e.g., eyes have been locally increased – see Figure 6). Several studies confirmed that increasing facial parts lowers perceived appeal, especially in case of realistic characters. In addition, Green et al. [2008] demonstrated that not only proportions, but also the placement of facial parts may affect negatively perceived appeal. The measured effect was bigger in cases where the original faces were more attractive and human-like. The results have been later confirmed [Burleigh et al. 2013; MacDorman et al. 2009] and partly extended by demonstrating that a mismatch of realism between facial parts negatively affects appeal.

The previous studies addressed the perception of rather unusual facial proportions for realistic characters and their influence on perceived appeal. In case of real people, previous work demonstrated that some facial proportions are associated with personality traits. A meta-analysis study concluded that individuals with wider faces were judged by observers as more threatening, more dominant and less attractive, especially for male faces [Geniole et al. 2015]. In addition larger eyes increase trustworthiness [Zebrowitz et al. 1996], while narrow eyes appear aggressively. Parts of these results have been confirmed for virtual characters. Narrow eyes have been rated as more aggressive and less trustworthy for both, abstract creatures [Ferstl et al. 2017] and more realistic virtual humans [Ferstl and McDonnell 2018]. It should be noted that eye size should not be modelled by varying the size of the eyes itself as this will be quickly perceived as eerie and artificial, but rather by changing the shape of the eyelids and partly the proportions of the head. Protruding eyes appear larger, whereas, hooded eyes and monolid eyes appear smaller. For virtual characters, the opposite result was found for the perception of wide faces, which were perceived as less aggressive and dominant [Wang et al. 2013] even when a masculine rather than a babyface appearance was achieved [Ferstl and McDonnell 2018].

5.4 Level of Detail

Creating models with different level of detail (LOD) is especially common in real-time applications, crowd scenes and detailed scenes like cities, terrains etc. The overall goal is to maintain rendering speed and a small memory footprint without loosing visual accuracy. Perceptual research on LOD can be divided in two categories: When do humans notice the artifacts of low-quality models and how are low quality models perceived when their smaller resolution is obvious.

Luebke et. al. [2003] provide probably one of the most detailed descriptions on perception within the LOD context and for topics related to crowds we refer to dedicated tutorials [Badler et al. 2014; Donikian et al. 2009; Huerre et al. 2010; Thalmann et al. 2006, 2007]. Within the context of virtual characters, two different representations exist to represent characters at lower resolutions. One option is to use textured, low resolution meshes. The other option are impostor techniques, where the object becomes a plane with albedo textures, transparency and normal maps. To maintain the 3d illusion, textures are rendered from different views and are replaced later depending on the user perspective. For computer displays, impostor representations remain indistinguishable until a pixel-to-texel ratio of slightly above one-to-one [Hamill et al. 2005]. Flickering artifacts become visible when changing the representation from impostor to 3d mesh at a pixel-to-texel ratio of one-to-one [McDonnell et al. 2006]. Furthermore, impostors are better at reproducing fine scale deformation and subtle motion than low resolution models. Differences in the view direction of $10-20^\circ$ remain unnoticeable. It should be noted that the authors of these studies mentioned visual artifacts due to aliasing as a reason for identification of impostors or transitions. Given the strong improvements in anti-aliasing algorithms, both within the rendered image as well as over the temporal domain, the low quality model may remain longer indistinguishable within current game engines. In the case of low-resolution models, visible artifacts can be identified due to a lack of smoothness within the silhouettes, incorrect lighting and texture distortion. Studying each factor independently for virtual characters, Larkin and O'Sullivan [2011] showed that silhouette is the dominant artifact for simplification identification at smaller screen spaces and lighting and silhouette artifacts are easily detected at larger screen spaces. However, when using normal maps, lighting artifacts can be masked efficiently.

Perception of low quality models when their smaller resolution is obvious is a side track of realistic vs. non-realistic character perception research. MacDorman et al. [2009] showed participants several images of virtual faces, combining different textures (from realistic to simple lines) with geometric levels of detail, where geometric detail was defined by the polygon count. Results suggested that reducing photo-realism can make the face look less eerie and more attractive. Similarly, Burleigh et al. [2013] compared faces with enlarged lips and eyes of different mesh resolutions. Faces

Figure 7: *Left to right:* Original realistic character, feature preserving texture blurring, stylized texture, blurred 2D rendering. Feature preserving blurring of realistic textures (not the image) increases appeal and is at the same time an intermediate step towards stylized textures.

with the lowest mesh resolution had lower variation on the perceived appeal when comparing normal and increased eye size (see Section 5.3). Furthermore, we also want to point the reader to the general role of perception within the context of mesh compression, where perceptual metrics are increasingly used to control the lossy compression locally [Corsini et al. 2012].

5.5 Skin Appearance

Besides dedicated work on skin appearance of virtual characters, relevant research has been carried out in the cosmetics research as well as in general research on attractiveness of people. Many studies concerning attractiveness of human faces merged different photographs to achieve average appearance. There was speculation that this technique impacts ratings of attractiveness not just because it averages the shape, but also because it removes blemishes and other skin irregularities [Alley and Cunningham 1991]. Several studies confirmed that texture changes do result in a significantly more attractive face [Benson and Perrett 1992; Little and Hancock 2002]. In the cosmetics domain, Fink et al. [2006] created textures from photographs of women of different ages and evaluated these textures on a single female virtual character. Renderings with pure skin were rated as younger and more attractive than renderings with strong variations in skin pigmentation. This observation was confirmed in a follow-up study, which showed that blurring the skin texture can increase attractiveness [Fink and Matts 2008]. Zell et al. [2015] observed similar effects for stylized characters as well. Blurring realistic textures, while preserving feature contours (e.g., lip contours) made characters with realistic shape as well as stylized shapes more appealing. In fact, texture stylization can be considered as a process that makes textures more uniform up to the point when features disappear depending on their visual importance (Figure 7). Empirical observations that smoother skin is considered more appealing can also be found in many photograph retouching books (e.g., [Nitzsche and Rose 2011]) and photo-retouching software for faces.

5.6 Lighting and Shading

The number of papers addressing different effects of lighting and shading within the context of virtual characters is small, but we can gain valuable knowledge by taking into account general studies

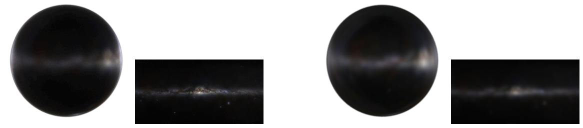

(a) Rough (b) Environment (c) Chrome (d) Environment

Figure 8: A special case demonstrating the ambiguity of light and material. *Left:* A sphere with a rough surface lit by a star environment map. *Right:* A chromatic sphere lit by a blurred environment map. Despite different materials and lighting, the appearance of both spheres is nearly identical.

on lighting and shading. We focus at this point on glossy, diffuse and translucent surfaces and omit transparent materials as these are less relevant in the context of virtual characters. Studying the perception of materials is challenging due to the strong interaction with lighting conditions. Certain visual appearances can be achieved by either modifying the surface appearance or the environment lighting as demonstrated in the extreme example in (Figure 8). Consider, for example, a perfectly polished chrome ball within a closed box covered with velvet. In the presence of indirect light, the ball would mirror the velvet surface, making it impossible to distinguish it from the velvet material of the box. It is therefore not surprising that participants inconsistently approximate parameters for glossy surfaces, especially in case of unnatural lighting [Fleming et al. 2003]. Also, accuracy of identifying equal materials [Pont and te Pas 2006] or the determination of roughness [Ho et al. 2006] vary for different lighting setups. Besides light, even the shape of an object influences the perception of glossiness [Olkkonen and Brainard 2010, 2011; Vangorp 2009; Vangorp et al. 2007]. Depending whether a small or big fraction of the surface area is covered by highlights, the material will be perceived as more or less glossy. Such ambiguity can only be resolved by providing several views of the same object. At the same time, the human visual system has developed an incredible ability to account for contextual information as well as surface properties in order to preserve the identity of an observed object. For instance, a black box remains identifiable as black no matter how bright the light within the room is. While humans perceive the box as black under different lighting conditions, the color will range between different shades of dark grey. This adaptation to contextual information is referred to as lightness constancy and it is a major challenge in visual science [Brainard 2003; Gilchrist et al. 1999].

In the case of translucent materials, the lighting direction has a fundamental impact on perceived translucency. While frontally lit translucent objects lack many visual cues (e.g., blurred features, soft shadows, low contrast), these features are enhanced when illuminated from the back [Fleming and Bülthoff 2005]. This effect is exploited in skin and hair rendering, where accurate shading models that replicate the physical behaviour of light are rendered with a back-light to visualize the fidelity of the shading model. In contrast, shaders that focus on performance and sacrifice accuracy are shown under less extreme lighting setups to underline their visual equivalence. The fact that the human visual system is tolerant to inaccuracies in lighting or shading was considered to speed-up rendering, e.g., approximating indirect light between

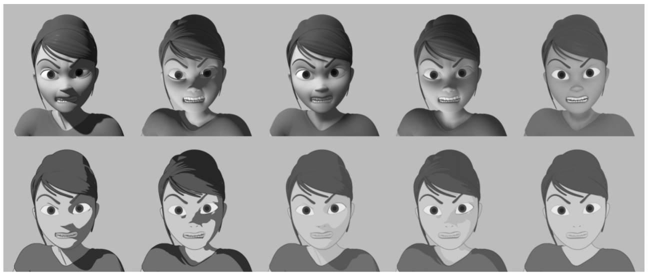

Figure 9: Stimuli from Wisessing et al. [2016] for testing two rendering styles (top/bottom) and five lighting conditions. From left to right: High contrast/light from above, High contrast/light from below, Low contrast/light from above, Low contrast/light from below, No directional Light.

frames through spherical harmonics [Jarabo et al. 2012]. Another well-established example is the replacement of the computationally intensive inter-reflections between surfaces through several simpler light sources [Keller 1997]. Respective perception parameters have been systematically studied in [Křivánek et al. 2010].

Rather than focusing on accuracy in material perception, the question remains what makes surfaces and light look realistic. Based on real photographs, Rademacher et al. [2001] identified that surface smoothness and shadow softness increase realism, but not the number of objects and lights. However, the effect size was bigger for surface smoothness than for light [Rademacher 2002]. A similar approach is considered as good practice among digital artists, who create multi-layered materials, with dedicated textures for surface scratches, dirt etc. It should be noted that it is often sufficient to add plausible dirt textures, but not necessarily replicate exactly the dirt.

Zell et al. [2015], investigated the perceived differences between accurate lighting and shading models across different stylization levels of virtual characters. Participants judged a character with Phong shading lit by simple directional lights with hard shadows almost as realistic as the same character with complex materials in combination with global illumination. While this result might sound surprising within the computer graphics community, Kardos [1934] mentioned 80 years ago that people tend to ignore shadows and shading when describing a scene.

Finally, lighting is considered as a powerful tool in cinematography to emphasize the mood within a scene. So far the majority of conclusions are mainly drawn based on observations [Grodal 2005]. The number of empirical studies testing the conclusions is rather small and does not always align with film theory. For example, a recent empirical study by Poland [2015] found that low-key/high contrast stimuli produced lightheartedness, contrary to the beliefs of many theorists and cinematographers. Within the context of virtual characters, Wisessing et al. [2016] (Figure 9) measured the impact of render style and lighting on the intensity and appeal of expressions in short animation sequences. Different lighting directions, such as the key light placed above or below the character had very little influence on perceived emotional intensity, and dark shadows were rated low on appeal.

5.7 Visual Attention of Facial and Body parts

Besides focusing on different aspects that contribute towards the appearance of a character, one should also consider that some body parts are more important than others. By using eye-tracking, McDonnell et al. [2009c] identified that viewers look mainly at the head and upper torso and used this information to create diverse looking crowds more effectively. In the case of faces, it is has been known for a long time [Groner et al. 1984] that people are looking primarily at the eyes and mouth but the number of fixations at the eyes dominate, which was later confirmed for realistic renderings [Raiturkar et al. 2018]. A recent study [Schwind and Jäger 2016] showed that this is also true for virtual characters of different stylization levels. On average, participants looked for 35% of the time at the eyes, while other regions ranged between 0-10%. This may explain why eyes are considered by practitioners as the most important aspect to achieve realism.

6 CHARACTER MOTION & BEHAVIOUR

Motion or animation of virtual characters is an integral part of the character design and can be achieved through artist animation (key-framing extreme poses and adding in-between frames), physics based animation (computer generated motion based on physics laws), and retrieving the actual motion from real life and applying it to a character (rotoscoping, depth cameras, motion capture), and combinations of those (synthesized motion). Animation approaches depend on the types of use, and they each have perceptually based rules which determine their success in creating a convincing character. In this course, we focus on the perception of realistic motion retrieved from motion capture and retargeted onto a virtual character; we are not interested in all the possible combinations, such as synthesized or procedural motions used in interactive scenarios, since the description of all approaches would result in a lengthy analysis. We therefore leave out some interactive components of virtual characters (see for example approaches to generate artificial eye-gaze behaviour [Ruhland et al. 2015a]).

People use non-verbal signs of the character such as motion and appearance to formulate opinions, judgments, or feelings about the character. A lot of social information is expressed through motion. Early work on the perception of biological human motion was done by putting lights on parts of the human body in a darkened room, the so called *point-light displays.* When the human was static, all that was perceived was a group of dots. When moving, the viewers could identify a human body engaged in a readily identified activity, such as walking, running, or dancing [Johansson 1973]. When studying biological motion applied to virtual characters, both shape and motion information interact to formulate a perceptual effect. It was found that the detection of biological motion can be obscured by increasing the anthropomorphism of the character [Chaminade et al. 2007].

Motion can also carry information about gender, emotion and personality of the mover. These motion specifics become very important when building virtual characters which are animated using natural motion from motion capture. In the next sections, we expand on these three types of information coming from motion and explain how they affect the perception of characters of different appearance.

6.1 Gender

The research using point-light displays have shown that gender can be recognized from motion when very little shape information is present. Men and women have a specific way of walking, and these differences are apparent to the observers: a pronounced sway in the area of hips often indicates a female walker, while a defined movement in the shoulder area indicates that it is a male walker [Kozlowski and Cutting 1977]. Not only walking, but also conversational motions (hand gestures, posture) applied to male, female and androgynous characters, can be accurately recognized as male or female motions [McDonnell and O'Sullivan 2010]. The participants in this study reported focusing on pose and wrist motions in order to estimate the underlying gender of the mover. Gender can be recognized from facial motions as well [Hill and Johnston 2001], where females can be discerned from males primarily because of more frequent nodding, blinking and overall amount of movement [Morrison et al. 2007].

Character appearance can affect the perception of gender from motion as well. For motions sparse on gender cues, it is the appearance of the character that will dominate our perception of the character's gender, whereas it is the motion that dominates the perception of characters with an androgynous appearance [McDonnell et al. 2009a]. In the case where motion with strong gender cues is applied to a virtual character of the opposite sex, e.g., male walking motion on a female character, it could result in the "contrast effect". Due to this effect, the gender from motion will be perceived even stronger when there is a mismatch with the gender of the character, much like a white paper will appear even whiter when put on a black background. Therefore, a male motion applied to a female character (and vice versa) may actually seem more "manly" due to such contrast [Zibrek et al. 2015]. Interestingly, this effect was dependent on the gender of the observer - males could identify female motions better and females could identify male motions better on a character of a mismatched sex. This example could point to a selective sensitivity when perceiving gender, perhaps due to the evolutionary importance in correctly recognizing the opposite sex, but also shows the importance of controlling for gender of the participant when conducting perceptual studies.

6.2 Emotion

Because the perception and interpretation of other people's emotion is essential for effective social interaction, people will find the character more engaging when it accurately expresses emotions. And since we put so much importance on emotions in our everyday life, the ability to recognize and distinguish between different emotional states has a prominent role in perceptual processes. Studying the perception of emotion is challenging, as there have been many attempts to define emotions and the exploration of their origin and development is an ongoing research focus [Lewis et al. 2008]. The most general definition describes emotions as subjective experiences, where the core feeling is that of pleasure or pain [Frijda 1988]. Several approaches to emotion classification exist in the literature as well, from defining emotion as discrete categories [Ekman 1992] or as dimensions [Mehrabian 1980; Plutchik 2001]. Ekman's approach to emotions was aimed at identifying emotions which are universally recognized and where similarity in their physical expression can be observed. He classified them as basic emotions: anger, happiness, sadness, fear, surprise and disgust. This classification also provided a simple as well as systematical approach for the study of emotion recognition from motion, which provided a comprehensive way to map emotions onto virtual characters. For example, in facial animation, the classification known as the Facial Action Coding System (FACS) is used for the creation of blend-shapes. Dimensional approach to classifying emotions has shown practical use as well, where Russell's circumplex model of emotion [Russell 1980], describing emotions in terms of valence (positive, negative) and activation (activation, deactivation) led to synthesized motion generation of complex facial expressions for characters [Grammer and Oberzaucher 2006].

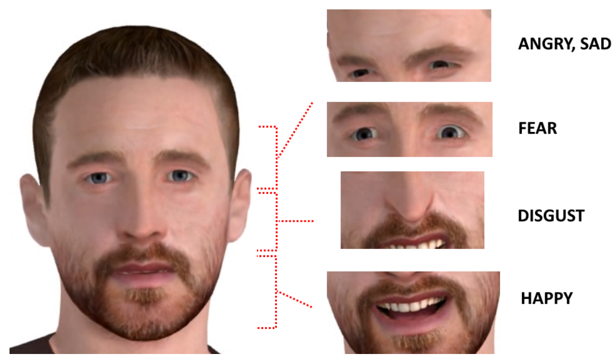

Figure 10: Areas of the face that carry the most information about emotions: mouth in happy, eyes in fear, disgust in nose and upper mouth areas, while sad and angry are expressed mostly with brow and eyes.

Basic emotions can be identified through the movement of full body [Atkinson et al. 2004; Crane and Gross 2007], upper body [Volkova et al. 2014] and arm motions [Pollick et al. 2001], and best emotion recognition rates are achieved when facial motions are combined [Clavel et al. 2009; Ennis et al. 2013; Hodgins et al. 2010], showing that a virtual character will be engaging when both face and body correctly express emotions in the animation. The accuracy of facial expressions is particularly important, as shown by the study of Hodgins et al. [2010], where facial motion anomalies were particularly salient even when obvious body motion anomalies were present. Research also shows that there are areas of the face which are important for particular emotions [Cunningham et al. 2004; Schyns et al. 2009]: happiness and surprise are expressed mostly with the mouth; sadness, anger and fear with the eyes and brows, as seen in Figure 10. One particular study integrated this knowledge to reverse-engineer facial expressions and improve recognition of emotions in faces of social robots [Chen et al. 2018]. In terms of body expressions [De Meijer 1989], the trunk was found to be the most important in conveying positive emotions. In the study of Zibrek et al. [2015], happiness and anger were found to be better expressed with facial and hand movements, while sad and fearful emotions were more apparent in full body motions.

Biological motion can therefore be used to create identifiable emotions for virtual characters. Emotion recognition from body motion is quite robust across different styles of character models,

Figure 11: Emotions were recognized equally well on the different bodies investigated by [McDonnell et al. 2009b]: androgynous mannequin, cute character, zombie, and point light display.

as the study of McDonnell [2009b] using a range of virtual characters (Figure 11), from point-light displays to high-fidelity shapes. However, sometimes character's appearance reduces the emotion recognition. For example, more texture information of the skin and some very abstract styles of rendering which introduce a lot of details to the face, interfere with efficient emotion perception from moving faces [Wallraven et al. 2007] or dampen the expression intensity [Hyde et al. 2013]. While visual fidelity of the character also requires complex texture information, which can dampen the emotion expression, it was found that in behavioral scenarios, it can also change the emotional experience of the viewer. In the study of Volante et al. [2016], it was shown that a realistically rendered patient, whose health is slowly deteriorating, unexpectedly increased the feelings of shame and shyness in medical students. The cartoon and sketch rendered styles had a higher value of the expected negative emotions and were more appropriate for inducing a stronger empathetic response, however, the authors suggest that a realistic character added nuance and complexity to the response of the students, making the experience more comparable to a real life situation. A similar finding was reported by Zibrek et al. [2018] where viewers were more concerned about the realistically rendered character when he was expressing anger and frustration than other styles (toon CG, toon shaded, zombie, creepy).

6.3 Gender and Emotion bias

When investigating virtual characters, some perceptual effects which are related to our experience in social interactions may arise. Particularly gender stereotypes, which are a consequence of cultural conditioning, impact both the production and perception of emotions [Brewer 1988]. Studies done with subjects living in Western societies show that overall, emotions are perceived to be gender specific. Females were found to be generally more expressive than males (and better at recognizing emotions of others as well) [Battocchi et al. 2005] but certain emotions were more likely to be expressed by males, e.g., anger, contempt and pride [Plant et al. 2000]. Sadness and fear were found to be more readily expressed by females [Fischer et al. 2004].

Based on this knowledge from investigating gender differences in perception and generation of emotion, studies investigated whether virtual characters would be subjected to bias as well. The study of Johnson et al. [2011] explored sex recognition bias on the perception of throwing a ball with an emotional motion style using point-light displays. They found that an angry throw is perceived as more male and a sad throw as more female, which supports the view that anger is more readily attributed to males than females. This view was extended to full body and conversational motions in the study of Zibrek et al. [2015]. Here, gender bias was explored on different types of motion with obvious gender cues (walking) and less obvious cues (conversation), while the motions were applied on male and female virtual characters. They found that emotion biases gender perception according to gender stereotypes: an angry motion is seen as more male, while fear and sadness are seen as less male motions. These studies show that perception of motion is influenced by the type of expression - full body motion has more gender cues than facial or hand motion alone, and emotional expression or appearance of the character will not influence the perception of gender. However, when motion type does not give enough information about the underlying gender, appearance and emotion will both influence the way the biological motion is perceived (see Figure 12).

Figure 12: Examples of emotion expressions from conversational motions (*top*) and walking motions (*bottom*) from Zibrek et al. [2015]. Motions from both male and female actors were applied to virtual characters of both sexes to explore the perceptual bias.

6.4 Personality

In order to create a more complex and engaging behaviour of virtual characters, one can consider designing them to express personality traits. As with emotion, different personality theories exist in the literature. Due to its continuous examination and re-evaluation, the "Big Five" theory [Costa and McCrae 1992; Goldberg 1990; John et al. 2008] is perceived by many to be the standard description of human personality. The Big Five is a hierarchical model of personality traits with five broad factors (extraversion, agreeableness, conscientiousness, openness to experience, and emotional stability). Each factor is bipolar (e.g., extraversion vs. introversion) and is further described by specific facets and traits. For example, extraverts are talkative and sociable, whereas introverts like to keep to themselves. Emotionally stable people do not get upset easily and keep a calm attitude even in stressful situations, while neurotic people will be more easily unnerved. People often assess other people on very little information and make judgments about their personality without knowing them before (see studies of zero-acquaintance [Borkenau and Liebler 1992; Mehl et al. 2006]): conscientiousness is associated with people who are well groomed and dress smart, while an extraverted person could be identified by having a very expressive face and hand gestures, and just being overall physically animated and energetic. A non-agreeable person could be identified by his or her way of conversing with others - tends to be quarrelsome, callous.

Some of the indicators of particular personality traits are visual (appearance, motion), while others are more associated with language and interactive behaviour. When this information is used to design virtual characters, perception can be steered towards a belief that the character possesses personality traits. It was found that exaggerations in body motion, the general speed of body motion [Neff et al. 2010], as well as increased speed of facial motion [Hyde et al. 2013] create a perception of the character as being more extraverted. A recent study on body shape created by using the skinned multi-person linear model found that for male bodies, extraversion was inferred from their wide shoulders and triangular body shape, while non-agreeable and neurotic personality types in women were associated with heavy-bottom, short-legged powerful body shapes [Hu et al. 2018]. Some inadequacies in a character's motion can also transfer to the observed personality disorders of the character. For example, Tinwell [2013] found that virtual characters with inadequate upper facial animation exhibit personality traits associated with psychopathy. Gaze behaviour was found to be very closely linked to personality as well [Rauthmann et al. 2012; Ruhland et al. 2015b], where characters who avert eye-contact in a conversation will appear more neurotic/introverted and long, uninterrupted eye-contact a sign of extraversion [Ruhland et al. 2015b].

It is possible therefore to create distinctive personality traits and these can be systematically investigated. In the study of Zibrek et al. [2014] for example, the Big Five traits were found to be distinguished by the participants and have a different effect in combination with the render style they were rendered in. In sum, there are many indications that appearance influences the perceived personality, as observed in real people. This creates a fascinating new venue for researchers in character perception to explore.

6.5 Distinctiveness

Individuals move in different ways, and this distinctiveness in motion is perceived by the observers. Distinctiveness of motion style can be observed in differences between examples of the same behaviour (e.g., slow walk vs. fast walk). When designing virtual humans, human motion style was studied to create motion models, which can be used to add subtle differences to the characters' motion when rendering virtual crowds [Ma et al. 2010] or for creating synthetic motion or transitions in motion captured motion [Lee et al. 2002]. However, no complete categorisation of motion styles, which would help with the creation of complex motion models, exists yet.

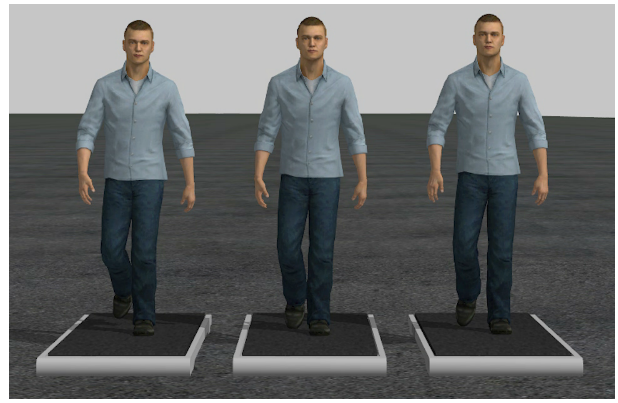

Figure 13: Example stimuli from [Hoyet et al. 2013], where they compared motions of different walkers for distinctiveness.

Distinctiveness also helps to identify familiar people based on motion only, which means that motion carries identity cues as well [Cutting and Kozlowski 1977]. However, this performance is quite poor: in the study of Cutting and Kozlowski [1977] only 38% of motions were correctly identified, where 16.7% was the chance level of guessing the correct identity. It was also found that people have difficulty distinguishing walking motions of different people [McDonnell et al. 2008] and it was found that only three distinct types of walking motions are necessary to create variety for virtual walkers in a virtual crowd [Pražák and O'Sullivan 2011].

The research on distinctiveness of body motion found a negative relationship with attractiveness - less distinctive movement, created by averaging similar types of motion (walking, jogging, dance) is perceived as more attractive [Hoyet et al. 2013] (Figure 13). The reason for that is perhaps in the lack of distinctiveness (averageness), which was shown to play an important role in perceived attractiveness of real faces [Rhodes 2006]. However, a particular type of distinctiveness, with sexually exaggerated cues, e.g., feminine movement in women characters, has been shown to be attractive as well [Johnson and Tassinary 2007].

7 LIMITATIONS

Our intention within this course was to focus primarily on the aspects in virtual character perception that are considered most relevant within the computer graphics domain, namely character appearance and motion. For this reason, we did not cover other modalities such as audio, touch, etc. Additionally, virtual character perception is covered within different domains, such as psychology, computer graphics, robotics, virtual reality and computer vision, making it difficult to identify relevant work. While we have tried to include research from different domains, we cannot guarantee to cover all relevant publications.

8 OPEN ISSUES

Despite the number of publications on the topic, we are still far from fully understanding the perception of virtual characters. Virtual characters have been present in various domains for many years. Today's quick developments of real-time rendering technologies enable communication in virtual environments (VR, AR), and social media. The need for understanding the implications of different forms of self-representations of the user either as a doppelgänger or an artificial avatar has never been greater.

Moreover, a more frequent exposure to virtual characters may continuously change the way we perceive them, much like how we are becoming more and more sensitive to bad visual effects in movies. It is relevant to note, that the perception of virtual characters is not a static but quite a dynamic research area which needs continuous reassessment.

One particular problem of virtual characters is their complexity, both in terms of their creation as well as modelling higher-level features to represent consistent behaviour. This multidimensionality has the drawback that some factors might cancel out or amplify each other, leading to inconsistent conclusions. For this reason, we recommend to begin with the investigation of one or two variables alone and increase complexity later. In addition, while complexity of stimuli and variety of implementations prevents a unified approach to studying virtual characters, we should strive to develop more standardized measures and unified terminology in the research field to increase the validity of our results.

One often encountered problem of perceptual studies is the access to high quality data and algorithms. A big thank you to all researchers and artists who made their assets accessible. Publication of assets and databases do not only allow access to state-of-the-art research, but add a baseline for other studies.

Today, we know that the eyes and mouth contribute strongly to the realism of a virtual character, but we do not know exactly what is required to make the characters look more realistic. Focusing more on salient facial and body parts in combination with stronger consideration of low-level perception could offer a clearer understanding of the important features. This could not only help to blur further the line between computer generated imaginary and photo-realism but offers interesting insights for designing algorithms to build digital doubles.

Finally, it should be emphasized that many interesting discoveries in character perception were made by rather unconventional experiment designs (e.g., point light displays, local scaling of facial parts). We encourage authors to consider prototyping a "crazy" experiment before dismissing the idea.

ACKNOWLEDGEMENTS

This research was funded by the Science Foundation Ireland as part of the "Game Face" project (grant 13/CDA/2135) and ADAPT Centre for Digital Content Technology (grant 13/RC/2106) projects. We would like to thank Mario Botsch, Thomas Waltemate, Martin Breidt, and Ludovic Hoyet for giving permission to add pictures of their work to this course.

REFERENCES

Sun Joo Ahn and Jeremy N. Bailenson. 2011. Self-endorsing versus other-endorsing in virtual environments: The effect on brand preference. *Journal of Advertising* (2011), 10–2753.

Thomas R. Alley and Michael R. Cunningham. 1991. Averaged faces are attractive, but very attractive faces are not average. *Psychological Science* 2, 2 (1991), 123–125.

Anthony P Atkinson, Winand H Dittrich, Andrew J Gemmell, Andrew W Young, et al. 2004. Emotion perception from dynamic and static body expressions in point-light and full-light displays. *Perception* 33, 6 (2004), 717–746.

Laura Aymerich-Franch, René F. Kizilcec, and Jeremy N. Bailenson. 2014. The Relationship between Virtual Self Similarity and Social Anxiety. *Frontiers in Human Neuroscience* 8 (2014), 944. https://doi.org/10.3389/fnhum.2014.00944

Norma Badler, Mubbasir Kapadia, Jan M. Allbeck, Yiorgos L. Chrysanthou, Nuria Pelechano, and Stephen Guy. 2014. Simulating heterogeneous crowds with interactive behaviors. In *Eurographics Tutorials*.

Jeremy N Bailenson, Kimberly R Swinth, Crystal L Hoyt, Susan Persky, Alex Dimov, and Jim Blascovich. 2005. The independent and interactive effects of embodied-agent appearance and behavior on self-report, cognitive, and behavioral markers of copresence in immersive virtual environments. *Presence* 14, 4 (2005), 379–393.

Domna Banakou, Raphaela Groten, and Mel Slater. 2013. Illusory ownership of a virtual child body causes overestimation of object sizes and implicit attitude changes. *Proc. of the National Academy of Sciences* 110, 31 (2013), 12846–12851.

Christoph Bartneck, Dana Kulić, Elizabeth Croft, and Susana Zoghbi. 2009. Measurement instruments for the anthropomorphism, animacy, likeability, perceived intelligence, and perceived safety of robots. *International journal of social robotics* 1, 1 (2009), 71–81.

Dirk Bartz, Douglas Cunningham, Jan Fischer, and Christian Wallraven. 2008. The role of perception for computer graphics. *Eurographics state-of-the-art-reports* (2008), 65–86.

Alberto Battocchi, Fabio Pianesi, and Dina Goren-Bar. 2005. A first evaluation study of a database of kinetic facial expressions (DaFEx). In *Proc. of Int. Conf. on Multimodal Interfaces (ICMI '05)*. 214–221.

David C. Beardslee and Michael Ed Wertheimer. 1958. Readings in perception. (1958).

Philip Benson and David Perrett. 1992. Face to face with the perfect image. *New Scientist* 133 (1992), 1809.

H. Richard Blackwell. 1952. Studies of psychophysical methods for measuring visual thresholds. *JOSA* 42, 9 (1952), 606–614.

Peter Borkenau and Anette Liebler. 1992. Trait inferences: Sources of validity at zero acquaintance. *Journal of Personality and Social Psychology* 62, 4 (1992), 645.

David H. Brainard. 2003. Color appearance and color difference specification. *The Science of Color* 2 (2003), 191–216.

Marilynn B. Brewer. 1988. *A dual process model of impression formation.* Lawrence Erlbaum Associates, Inc.

Tyler J. Burleigh, Jordan R. Schoenherr, and Guy L. Lacroix. 2013. Does the Uncanny Valley exist? An empirical test of the relationship between eeriness and the human likeness of digitally created faces. *Computers in Human Behavior* 29, 3 (2013), 759–771.

Thierry Chaminade, Jessica Hodgins, and Mitsuo Kawato. 2007. Anthropomorphism influences perception of computer-animated characters' actions. *Social Cognitive and Affective Neuroscience* 2, 3 (2007).

Marcus Cheetham, Pascal Suter, and Lutz Jäncke. 2011. The human likeness dimension of the Uncanny Valley hypothesis: Behavioral and functional MRI findings. *Frontiers in Human Neuroscience* 5 (2011), 126.

Chaona Chen, Oliver GB Garrod, Jiayu Zhan, Jonas Beskow, Philippe G Schyns, and Rachael E Jack. 2018. Reverse Engineering Psychologically Valid Facial Expressions of Emotion into Social Robots. In *Int. Conf. on Automatic Face & Gesture Recognition*. 448–452.

Céline Clavel, Justine Plessier, Jean-Claude Martin, Laurent Ach, and Benoit Morel. 2009. Combining Facial and Postural Expressions of Emotions in a Virtual Character. In *Proc. of Int. Conf. on Intelligent Virtual Agents (IVA '09)*. 287–300.

Massimiliano Corsini, Mohamed-Chaker Larabi, Guillaume Lavoué, Oldřich Petřík, Libor Vása, and Kai Wang. 2012. Perceptual Metrics for Static and Dynamic Triangle Meshes. In *Eurographics STAR*. 135–157.

Paul T. Costa and Robert R. McCrae. 1992. Professional manual: revised NEO personality inventory (NEO-PI-R) and NEO five-factor inventory (NEO-FFI). *Odessa, FL: Psychological Assessment Resources* (1992).

Elizabeth Crane and Melissa Gross. 2007. Motion capture and emotion: Affect detection in whole body movement. In *Affective computing and intelligent interaction*. Springer, 95–101.

Douglas P. Crowne and David Marlowe. 1960. A new scale of social desirability independent of psychopathology. *Journal of consulting psychology* 24, 4 (1960), 349.

Douglas Cunningham and Christian Wallraven. 2013. Understanding and Designing Perceptual Experiments. In *Eurographics Tutorial*.

Douglas W Cunningham, Mario Kleiner, Heirich H Bülthoff, and Christian Wallraven. 2004. The components of conversational facial expressions. In *Proc. of Symp. on Applied Perception in Graphics and Visualization*. ACM, 143–150.

James E Cutting and Lynn T Kozlowski. 1977. Recognizing friends by their walk: Gait perception without familiarity cues. *Bulletin of the psychonomic society* 9, 5 (1977), 353–356.

Zhengyan Dai and Karl F. MacDorman. 2018. The doctor's digital double: how warmth, competence, and animation promote adherence intention. *PeerJ Computer Science* 4 (2018).

Marco De Meijer. 1989. The contribution of general features of body movement to the attribution of emotions. *Journal of Nonverbal behavior* 13, 4 (1989), 247–268.

Vanderson Dill, Laura Mattos Flach, Rafael Hocevar, Christian Lykawka, Soraia Raupp Musse, and Márcio Sarroglia Pinho. 2012. Evaluation of the Uncanny Valley in CG Characters. In *Proc. of Intelligent Virtual Agents*. 511–513.

Stéphane Donikian, Nadia Magnenat-Thalmann, Julien Pettré, and Daniel Thalmann. 2009. In *Eurographics Tutorials*.

Paul Ekman. 1992. An argument for basic emotions. *Cognition and Emotion* 6, 3-4 (1992), 169–200.

Cathy Ennis, Ludovic Hoyet, Arjan Egges, and Rachel McDonnell. 2013. Emotion Capture: Emotionally Expressive Characters for Games. In *Proc. of Motion on Games*. ACM, 53–60.

Shaojing Fan, Rangding Wang, Tian-Tsong Ng, Cheston Y.-C. Tan, Jonathan S. Herberg, and Bryan L. Koenig. 2014. Human Perception of Visual Realism for Photo and Computer-Generated Face Images. *ACM Transaction Applied Perception* 11, 2, Article 7 (July 2014), 21 pages. https://doi.org/10.1145/2620030

Ylva Ferstl, Elena Kokkinara, and Rachel McDonnell. 2017. Facial Features of Non-player Creatures Can Influence Moral Decisions in Video Games. *ACM Transaction on Applied Perception* 15, 1, Article 4 (2017), 12 pages. https://doi.org/10.1145/3129561

Ylva Ferstl and Rachel McDonnell. 2018. A Perceptual Study on the Manipulation of Facial Features for Trait Portrayal in Virtual Agents. In *Proc. of Int. Conf. on Intelligent Virtual Agents (IVA)*. 281–288. https://doi.org/10.1145/3267851.3267891

Bernhard Fink, Karl Grammer, and Paul J. Matts. 2006. Visible skin color distribution plays a role in the perception of age, attractiveness, and health in female faces. *Evolution and Human Behavior* 27, 6 (2006), 845–857.

Bernhard Fink and Paul J. Matts. 2008. The effects of skin colour distribution and topography cues on the perception of female facial age and health. *Journal of the European Academy of Dermatology and Venereology* 22, 4 (2008), 493–498.

Agneta H. Fischer, Patricia M. Rodriguez Mosquera, Annelies E.M. Van Vianen, and Antony S.R. Manstead. 2004. Gender and culture differences in emotion. *Emotion* 4, 1 (2004), 87.

Reuben Fleming, Betty J. Mohler, Javier Romero, Michael J. Black, and Martin Breidt. 2016. Appealing Female Avatars from 3D Body Scans: Perceptual Effects of Stylization. In *Int. Conf. on Computer Graphics Theory and Applications (GRAPP)*.

Roland W. Fleming and Heinrich H. Bülthoff. 2005. Low-Level Image Cues in the Perception of Translucent Materials. *ACM Transactions on Applied Perception* 2, 3 (2005), 346–382.

Roland W. Fleming, Ron O. Dror, and Edward H Adelson. 2003. Real-world illumination and the perception of surface reflectance properties. *Journal of Vision* 3, 6 (2003), 347–368.

Roland W. Fleming and Manish Singh. 2009. Visual Perception of 3D Shape. In *ACM SIGGRAPH Courses*. Article 24, 94 pages.

Jesse. Fox and Jeremy. N. Bailenson. 2009. Virtual self-modeling: The effects of vicarious reinforcement and identification on exercise behaviors. *Media Psychology* 12 (2009), 1–25.

Nico H. Frijda. 1988. The laws of emotion. *American psychologist* 43, 5 (1988), 349.

Shawn N. Geniole, Thomas F. Denson, Barnaby J. Dixson, Justin M. Carré, and Cheryl M. McCormick. 2015. Evidence from Meta-Analyses of the Facial Width-to-Height Ratio as an Evolved Cue of Threat. *PloS one* 10, 7 (2015), e0132726.

Erik Geslin. 2012. Method of Induction of Basic and Complex Emotions in Video Games and Virtual Environments. In *SIGGRAPH Asia 2012 Courses*. Article 6, 70 pages.

Simon Gibson and Roger J. Hubbold. 1997. Perceptually-driven radiosity. In *Computer Graphics Forum*, Vol. 16. Wiley Online Library, 129–141.

Alan Gilchrist, Christos Kossyfidis, Frederick Bonato, Tiziano Agostini, Joseph Cataliotti, Xiaojun Li, Branka Spehar, Vidal Annan, and Elias Economou. 1999. An anchoring theory of lightness perception. *Psychological Review* 106, 4 (1999).

Mashhuda Glencross, Alan G. Chalmers, Ming C. Lin, Ming C. Lin, Miguel A. Otaduy, and Diego Gutierrez. 2006a. Exploiting Perception in High-fidelity Virtual environments. In *ACM SIGGRAPH 2006 Courses*. Article 1.

Mashhuda Glencross, Alan G. Chalmers, Ming C. Lin, Ming C. Lin, Miguel A. Otaduy, and Diego Gutierrez. 2006b. Exploiting Perception in High-fidelity Virtual environments. In *ACM SIGGRAPH 2006 Courses (SIGGRAPH '06)*. Article 1. http://doi.acm.org/10.1145/1185657.1185814

Lewis R. Goldberg. 1990. An alternative "description of personality": the big-five factor structure. *Journal of personality and social psychology* 59, 6 (1990), 1216.

Geoffrey Gorisse, Olivier Christmann, Samory Houzangbe, and Simon Richir. 2019. From Robot to Virtual Doppelganger: Impact of Visual Fidelity of Avatars Controlled in Third-Person Perspective on Embodiment and Behavior in Immersive Virtual Environments. *Frontiers in Robotics and AI* 6 (2019), 8. https://doi.org/10.3389/frobt.2019.00008

Karl Grammer and Elisabeth Oberzaucher. 2006. The reconstruction of facial expressions in embodied systems. *ZiF: mitteilungen* 2 (2006), 14–31.

Robert D. Green, Karl F. MacDorman, Chin-Chang Ho, and Sandosh Vasudevan. 2008. Sensitivity to the proportions of faces that vary in human likeness. *Computers in Human Behavior* 24, 5 (2008), 2456–2474.

Anthony G. Greenwald, Debbie E. McGhee, and Jordan L.K. Schwartz. 1998. Measuring individual differences in implicit cognition: the implicit association test. *Journal of personality and social psychology* 74, 6 (1998), 1464.

Torben Grodal. 2005. *Film Lighting and Mood. Moving Image Theory: Ecological Considerations.*

Rudolf Groner, Franziska Walder, and Marina Groner. 1984. Looking at faces: local and global aspects of scanpaths. *Theoretical and applied aspects of eye movements research* (1984), 523–533.

John Hamill, Rachel McDonnell, Simon Dobbyn, and Carol O'Sullivan. 2005. Perceptual evaluation of impostor representations for virtual humans and buildings. In *Computer Graphics Forum*, Vol. 24. Wiley Online Library, 623–633.

David Hanson. 2005. Expanding the Aesthetics Possibilities for Humanlike Robots. In *Proc. of IEEE Humanoid Robotics Conf., Special Session on the Uncanny Valley.*

Aaron Hertzmann, Carol O'Sullivan, and Ken Perlin. 2009. Realistic human body movement for emotional expressiveness. In *ACM SIGGRAPH 2009 Courses*. 1–27.

Harold Hill and Alan Johnston. 2001. Categorizing sex and identity from the biological motion of faces. *Current biology* 11, 11 (2001), 880–885.

Chin-Chang Ho and Karl F. MacDorman. 2010. Revisiting the Uncanny Valley theory: Developing and validating an alternative to the Godspeed indices. *Computers in Human Behavior* 26, 6 (2010), 1508–1518.

Yun-Xian Ho, Michael S. Landy, and Laurence T. Maloney. 2006. How direction of illumination affects visually perceived surface roughness. *Journal of Vision* 6, 8 (2006), 634–648.

Jessica Hodgins, Sophie Jörg, Carol O'Sullivan, Sang Il Park, and Moshe Mahler. 2010. The saliency of anomalies in animated human characters. *ACM Transactions on Applied Perception* 7, 4 (2010).

Ludovic Hoyet, Kenneth Ryall, Katja Zibrek, Hwangpil Park, Jehee Lee, Jessica Hodgins, and Carol O'Sullivan. 2013. Evaluating the Distinctiveness and Attractiveness of Human Motions on Realistic Virtual Bodies. *ACM Transaction Graphics* 32, 6 (2013), 204:1–204:11.

Ying Hu, Connor J. Parde, Matthew Q. Hill, Naureen Mahmood, and Alice J. O'Toole. 2018. First Impressions of Personality Traits From Body Shapes. *Psychological Science* 29, 12 (2018), 1969–1983.

Stephanie Huerre, Jehee Lee, Ming Lin, and Carol O'Sullivan. 2010. Simulating Believable Crowd and Group Behaviors. In *ACM SIGGRAPH ASIA 2010 Courses*. Article 13, 92 pages.

Jennifer Hyde, Elizabeth J. Carter, Sara Kiesler, and Jessica K. Hodgins. 2013. Perceptual Effects of Damped and Exaggerated Facial Motion in Animated Characters. (2013).

Adrian Jarabo, Tom Van Eyck, Veronica Sundstedt, Kavita Bala, Diego Gutierrez, and Carol O'Sullivan. 2012. Crowd Light: Evaluating the Perceived Fidelity of Illuminated Dynamic Scenes. *Computer Graphics Forum* 31, 2 (2012), 565–574.

Arthur R. Jensen. 1965. Scoring the Stroop test. *Acta psychologica* 24, 5 (1965), 398–408.

Gunnar Johansson. 1973. Visual perception of biological motion and a model for its analysis. *Perception & Psychophysics* 14, 2 (1973), 201–211.

Oliver P John, Laura P Naumann, and Christopher J Soto. 2008. Paradigm shift to the integrative big five trait taxonomy. *Handbook of personality: Theory and research* 3 (2008), 114–158.

Kerri L. Johnson, Lawrie S. McKay, and Frank E. Pollick. 2011. He throws like a girl (but only when he's sad): Emotion affects sex-decoding of biological motion displays. *Cognition* 119, 2 (2011), 265–280.

Kerri L Johnson and Louis G Tassinary. 2007. Compatibility of basic social perceptions determines perceived attractiveness. *Proc. of the National Academy of Sciences* 104, 12 (2007), 5246–5251.

Yvonne Jung, Arjan Kuijper, Dieter W. Fellner, Michael Kipp, Jan Miksatko, Jonathan Gratch, and Daniel Thalmann. 2011. Believable Virtual Characters in Human-Computer Dialogs. In *Eurographics STAR*.

Ludwig Kardos. 1934. Ding und Schatten. Eine experimentelle Untersuchung Ãijber die Grundlagen des Farbensehens. *Zeitschrift für Psychologie und Physiologie der Sinnesorgane* (1934).

Jari Kätsyri, Klaus Förger, Meeri. Mäkäräinen, and Tapio Takala. 2015. A review of empirical evidence on different Uncanny Valley hypotheses: Support for perceptual mismatch as one road to the valley of eeriness. *Frontiers in Psychology* 6 (2015), 390.

Alexander Keller. 1997. Instant radiosity. In *Proc. of ACM SIGGRAPH*. 49–56.

Lynn T Kozlowski and James E Cutting. 1977. Recognizing the sex of a walker from a dynamic point-light display. *Perception & Psychophysics* 21, 6 (1977), 575–580.

Jaroslav Křivánek, James A. Ferwerda, and Kavita Bala. 2010. Effects of global illumination approximations on material appearance. *ACM Transaction on Graphics* 29, 4 (2010), 112:1–112:10.

Michéal Larkin and Carol O'Sullivan. 2011. Perception of Simplification Artifacts for Animated Characters. In *Proc. of Symp. on Applied Perception in Graphics and Visualization (APGV '11)*. 93–100. https://doi.org/10.1145/2077451.2077469

Jehee Lee, Jinxiang Chai, Paul SA Reitsma, Jessica K Hodgins, and Nancy S Pollard. 2002. Interactive control of avatars animated with human motion data. In *ACM Transaction on Graphics (TOG)*, Vol. 21. ACM, 491–500.

Michael Lewis, Jeannette M. Haviland-Jones, and Lisa F. Barrett. 2008. *Handbook of Emotions, Third Edition.* Guilford Publications.

Jamy Li, René Kizilcec, Jeremy Bailenson, and Wendy Ju. 2016. Social robots and virtual agents as lecturers for video instruction. *Computers in Human Behavior* 55 (2016), 1222 – 1230. https://doi.org/10.1016/j.chb.2015.04.005

Anthony C Little and Peter JB Hancock. 2002. The role of masculinity and distinctiveness in judgements of human male facial attractiveness. *British Journal of Psychology* 93, 4 (2002), 451–464.

Alexander D Logvinenko and Laurence T Maloney. 2006. The proximity structure of achromatic surface colors and the impossibility of asymmetric lightness matching. *Perception & Psychophysics* 68, 1 (2006), 76–83.

Christine E. Looser and Thalia Wheatley. 2010. The tipping point of animacy how, when, and where we perceive life in a face. *Psychological Science* 21, 12 (2010), 1854–1862.

David Luebke, Martin Reddy, Jonathan D. Cohen, Amitabh Varshney, Benjamin Watson, and Robert Huebner. 2003. *Level of Detail for 3D Graphics.* Morgan Kaufmann.

Wanli Ma, Shihong Xia, Jessica K Hodgins, Xiao Yang, Chunpeng Li, and Zhaoqi Wang. 2010. Modeling style and variation in human motion. In *Proc. of Symp. on Computer Animation.* Eurographics Association, 21–30.

Karl F. MacDorman. 2006. Subjective ratings of robot video clips for human likeness, familiarity, and eeriness: An exploration of the Uncanny Valley. In *Proc. of ICCS/CogSci Long Symp.: Towards Social Mechanisms of Android Science.*

Karl F. MacDorman, Robert D. Green, Chin-Chang Ho, and Clinton T. Koch. 2009. Too real for comfort? Uncanny responses to computer generated faces. *Computers in Human Behavior* 25, 3 (2009), 695–710.

Brandon Mader, Martin S. Banks, and Hany Farid. 2017. Identifying Computer-Generated Portraits: The Importance of Training and Incentives. *Perception* 46, 9 (2017), 1062–1076. https://doi.org/10.1177/0301006617713633

Katerina Mania and Erik Reinhard. 2008. Perceptually-Motivated Graphics. In *Eurographics Tutorials.*

Scott McCloud. 1993. *Understanding Comics.* William Morrow Paperbacks.

Rachel McDonnell, Martin Breidt, and Heinrich H. Bülthoff. 2012. Render me real? Investigating the effect of render style on the perception of animated virtual humans. *ACM Transaction on Graphics* 31, 4 (2012), 91:1–91:11.

Rachel McDonnell, Simon Dobbyn, Steven Collins, and Carol O'Sullivan. 2006. Perceptual Evaluation of LOD Clothing for Virtual Humans. In *Proc. of Symp. on Computer Animation (SCA '06).* 117–126.

Rachel McDonnell, Sophie Jörg, Jessica K Hodgins, Fiona Newell, and Carol O'sullivan. 2009a. Evaluating the effect of motion and body shape on the perceived sex of virtual characters. *ACM Transactions on Applied Perception (TAP)* 5, 4 (2009), 20.

Rachel McDonnell, Sophie Jörg, Joanna McHugh, Fiona N Newell, and Carol O'Sullivan. 2009b. Investigating the role of body shape on the perception of emotion. *ACM Transactions on Applied Perception (TAP)* 6, 3 (2009), 14.

Rachel McDonnell, Michéal Larkin, Simon Dobbyn, Steven Collins, and Carol O'Sullivan. 2008. Clone attack! perception of crowd variety. *ACM Transaction on Graphics (TOG)* 27, 3 (2008), 26.

Rachel McDonnell, Micheal Larkin, Benjamin Hernandez, Isaac Rudomin, , and Carol O'Sullivan. 2009c. Eye-catching Crowds: Saliency based Selective Variation. *ACM Transaction on Graphics* 28, 3 (2009), 55:1 – 55:10.

Rachel McDonnell and Carol O'Sullivan. 2010. Movements and voices affect perceived sex of virtual conversers. In *Proc. of Symp. on Applied Perception in Graphics and Visualization.* ACM, 125–128.

Ann McNamara, Katerina Mania, and Diego Gutierrez. 2011. Perception in Graphics, Visualization, Virtual Environments and Animation. In *SIGGRAPH Asia 2011 Courses.* Article 17, 137 pages.

Matthias R Mehl, Samuel D Gosling, and James W Pennebaker. 2006. Personality in its natural habitat: Manifestations and implicit folk theories of personality in daily life. *Journal of personality and social psychology* 90, 5 (2006), 862.

Albert Mehrabian. 1980. Basic Dimensions for a General Psychological Theory Implications for Personality, Social, Environmental, and Developmental Studies. (1980).

Edward R. Morrison, Lisa Gralewski, Neill Campbell, and Ian S. Penton-Voak. 2007. Facial movement varies by sex and is related to attractiveness. *Evolution and Human Behavior* 28, 3 (2007), 186–192.

Michael Neff, Yingying Wang, Rob Abbott, and Marilyn Walker. 2010. Evaluating the effect of gesture and language on personality perception in conversational agents. In *Int. Conf. on Intelligent Virtual Agents.* Springer, 222–235.

Paula M Niedenthal, Jamin B Halberstadt, Jonathan Margolin, and Åse H Innes-Ker. 2000. Emotional state and the detection of change in facial expression of emotion. *European Journal of Social Psychology* 30, 2 (2000), 211–222.

Birgit Nitzsche and Karsten Rose. 2011. *Bodyshop: The Photoshop Retouching Guide for the Face and Body.* Wiley.

Maria Olkkonen and David H. Brainard. 2010. Perceived glossiness and lightness under real-world illumination. *Journal of Vision* 10, 9 (2010).

Maria Olkkonen and David H. Brainard. 2011. Joint effects of illumination geometry and object shape in the perception of surface reflectance. *i-Perception* 2 (2011), 1014–1034.

Carol O'Sullivan, Sarah Howlett, Yann Morvan, Rachel McDonnell, and Keith O'Conor. 2004. Perceptually Adaptive Graphics. In *Eurographics STAR.*

Ye Pan and Anthony Steed. 2016. A Comparison of Avatar-, Video-, and Robot-Mediated Interaction on Users Trust in Expertise. *Frontiers in Robotics and AI* 3 (2016), 12. https://doi.org/10.3389/frobt.2016.00012

E Ashby Plant, Janet Shibley Hyde, Dacher Keltner, and Patricia G Devine. 2000. The gender stereotyping of emotions. *Psychology of Women Quarterly* 24, 1 (2000), 81–92.

Robert Plutchik. 2001. The Nature of Emotions Human emotions have deep evolutionary roots, a fact that may explain their complexity and provide tools for clinical practice. *American Scientist* 89, 4 (2001), 344–350.

Jennifer Lee Poland. 2015. *Light, Camera, Emotion! An Examination on Film Lighting and It's Impact on Audiences' Emotional Response.* Cleveland State University.

Frank E Pollick, Helena M Paterson, Armin Bruderlin, and Anthony J Sanford. 2001. Perceiving affect from arm movement. *Cognition* 82, 2 (2001), B51–B61.

Sylvia C. Pont and Susan F. te Pas. 2006. Material – illumination ambiguities and the perception of solid objects. *Perception* 35 (2006), 1331–1350.

Martin Pražák and Carol O'Sullivan. 2011. Perceiving human motion variety. In *Proc. of Symp. on Applied Perception in Graphics and Visualization.* ACM, 87–92.

Pablo Mauricio Rademacher. 2002. *Measuring the Perceived Visual Realism of Images.* PhD Thesis.

Pablo Mauricio Rademacher, Jed Lengyel, Edward Cutrell, and Turner Whitted. 2001. Measuring the Perception of Visual Realism in Images. In *Proc. of the Eurographics Symp. on Rendering (EGSR).* 235–247.

Pallavi Raiturkar, Hany Farid, and Eakta Jain. 2018. *Identifying Computer-Generated Portraits: an Eye Tracking Study.* Technical Report. University of Florida.

John F Rauthmann, Christian T Seubert, Pierre Sachse, and Marco R Furtner. 2012. Eyes as windows to the soul: gazing behavior is related to personality. *Journal of Research in Personality* 46, 2 (2012), 147–156.

Gillian Rhodes. 2006. The evolutionary psychology of facial beauty. *Annu. Rev. Psychol.* 57 (2006), 199–226.

Kiaran Ritchie, Jake. Callery, and Karim. Biri. 2005. *The Art of Rigging: A Definitive Guide to Character Technical Direction with Alias Maya.* Vol. 1. Alias.

Kerstin Ruhland, Christopher E Peters, Sean Andrist, Jeremy B Badler, Norman I Badler, Michael Gleicher, Bilge Mutlu, and Rachel McDonnell. 2015a. A review of eye gaze in virtual agents, social robotics and hci: Behaviour generation, user interaction and perception. In *Computer graphics forum*, Vol. 34. Wiley Online Library, 299–326.

Kerstin Ruhland, Katja Zibrek, and Rachel McDonnell. 2015b. Perception of personality through eye gaze of realistic and cartoon models. In *Proc. of Symp. on Applied Perception.* ACM, 19–23.

James A Russell. 1980. A circumplex model of affect. *Journal of personality and social psychology* 39, 6 (1980), 1161.

Edward Schneider, Yifan Wang, and Shanshan Yang. 2007. Exploring the Uncanny Valley with Japanese Video Game Characters. In *Situated Play: Proc. of Digital Games Research Association Conference.*

D. Samuel Schwarzkopf, Chen Song, and Geraint Rees. 2011. The surface area of human V1 predicts the subjective experience of object size. *Nature neuroscience* 14, 1 (2011), 28.

Valentin Schwind and Solveigh Jäger. 2016. The Uncanny Valley and the Importance of Eye Contact. *i-com* (2016), 93–104. Issue 15.

Philippe G Schyns, Lucy S Petro, and Marie L Smith. 2009. Transmission of facial expressions of emotion co-evolved with their efficient decoding in the brain: behavioral and brain evidence. *PLoS One* 4, 5 (2009), e5625.

Kathryn Y. Segovia and Jeremy N. Bailenson. 2012. Virtual imposters: Responses to avatars that do not look like their controllers. *Social Influence* 7, 4 (2012), 285–303. https://doi.org/10.1080/15534510.2012.670906

Jun'ichiro Seyama and Ruth S. Nagayama. 2007. The Uncanny Valley: Effect of Realism on the Impression of Artificial Human Faces. *Presence: Teleoperators and Virtual Environments* 16, 4 (2007), 337–351.

Anne Thaler, Ivelina Piryankova, Jeanine K. Stefanucci, Sergi Pujades, Stephan de la Rosa, Stephan Streuber, Javier Romero, Michael J. Black, and Betty J. Mohler. 2018. Visual Perception and Evaluation of Photo-Realistic Self-Avatars From 3D Body Scans in Males and Females. *Frontiers in ICT* 5 (Sept. 2018), 1–14.

Daniel Thalmann, Carol O'Sullivan, Pablo de Heras Ciechomski, and Simon Dobbyn. 2006. Populating Virtual Environments with Crowds. In *Eurographics Tutorials.*

Daniel Thalmann, Carol O'Sullivan, Barbara Yersin, Jonathan Maim, and Rachel McDonnell. 2007. Populating Virtual Environments with Crowds. In *Eurographics Tutorials.*

Angela Tinwell, Deborah Abdel Nabi, and John P Charlton. 2013. Perception of psychopathy and the Uncanny Valley in virtual characters. *Computers in Human Behavior* 29, 4 (2013), 1617–1625.

Peter Vangorp. 2009. *Human Visual Perception of Materials in Realistic Computer Graphics.* PhD Thesis, University Leuven.

Peter Vangorp, Jurgen Laurijssen, and Philip Dutre. 2007. The Influence of Shape on the Perception of Material Reflectance. *ACM Transaction on Graphics* 26, 3 (2007), 77:1–77:9.

Matias Volante, Sabarish V Babu, Himanshu Chaturvedi, Nathan Newsome, Elham Ebrahimi, Tania Roy, Shaundra B Daily, and Tracy Fasolino. 2016. Effects of virtual human appearance fidelity on emotion contagion in affective inter-personal simulations. *IEEE Transaction on Visualization and Computer Graphics* 22, 4 (2016), 1326–1335.

Ekaterina P. Volkova, Betty J. Mohler, Trevor J. Dodds, Joachim Tesch, and Heinrich H. Bülthoff. 2014. Emotion categorization of body expressions in narrative scenarios. *Frontiers in Psychology* 5 (2014).

Christian Wallraven, Martin Breidt, Douglas W. Cunningham, and Heinrich H. Bülthoff. 2008. Evaluating the perceptual realism of animated facial expressions. *ACM Transactions on Applied Perception* 4, 4 (2008), 4:1–4:20.

Christian Wallraven, Heinrich H. Bülthoff, Douglas W. Cunningham, Jan Fischer, and Dirk Bartz. 2007. Evaluation of real-world and computer-generated stylized facial expressions. *ACM Transactions on Applied Perception* 4, 3, Article 16 (2007).

Thomas Waltemate, Dominik Gall, Daniel Roth, Mario Botsch, and Marc Erich Latoschik. 2018. The Impact of Avatar Personalization and Immersion on Virtual Body Ownership, Presence, and Emotional Response. *IEEE Transactions on Visualization and Computer Graphics* 24, 4 (2018), 1643–1652. https://doi.org/10.1109/TVCG.2018.2794629

Yuqiong Wang, Joe Geigel, and Andrew Herbert. 2013. Reading Personality: Avatar vs. Human Faces. In *Humaine Association Conference on Affective Computing and Intelligent Interaction*. 479–484. https://doi.org/10.1109/ACII.2013.85

Max Wertheimer. 1923. A brief introduction to gestalt, identifying key theories and principles. *Psychol Forsch* 4 (1923), 301–350.

Pisut Wisessing, John Dingliana, and Rachel McDonnell. 2016. Perception of Lighting and Shading for Animated Virtual Characters. In *Proc. of ACM Symp. of Applied Perception (SAP)*. 25–29.

Nick Yee, Jeremy N Bailenson, and Kathryn Rickertsen. 2007. A Meta-analysis of the Impact of the Inclusion and Realism of Human-like Faces on User Experiences in Interfaces. In *Proc. of the SIGCHI Conference on Human Factors in Computing Systems (CHI '07)*. 1–10. https://doi.org/10.1145/1240624.1240626

Lesslie A. Zebrowitz, Luminita Voinescu, and Mary Ann Collins. 1996. "wideeyed" and "crooked-faced": Determinants of perceived and real honesty across the life span. *Personality and Social Psychology Bulletin* 12, 12 (1996), 1258–âĂŞ1269.

Eduard Zell, Carlos Aliaga, Adrian Jarabo, Katja Zibrek, Diego Gutierrez, Rachel McDonnell, and Mario Botsch. 2015. To Stylize or Not to Stylize?: The Effect of Shape and Material Stylization on the Perception of Computer-generated Faces. *ACM Transactions on Graphics* 34, 6, Article 184 (2015), 184:1–184:12 pages.

Katja Zibrek, Ludovic Hoyet, Kerstin Ruhland, and Rachel McDonnell. 2015. Exploring the effect of motion type and emotions on the perception of gender in virtual humans. *ACM Transactions on Applied Perception (TAP)* 12, 3 (2015), 11.

Katja Zibrek, Elena Kokkinara, and Rachel McDonnell. 2018. The Effect of Realistic Appearance of Virtual Characters in Immersive Environments-Does the Character's Personality Play a Role? *IEEE Transactions on Visualization and Computer Graphics* 24, 4 (2018), 1681–1690.

Katja Zibrek and Rachel McDonnell. 2014. Does render style affect perception of personality in virtual humans?. In *Proc. of ACM Symp. on Applied Perception (SAP)*. 111–115.